


Architecture & Production Design

FrontierAtlas Ingestion Pipeline — Phase VI design document

Stack: Python 3.11+ · asyncio + aiohttp · BeautifulSoup/lxml · feedparser · Playwright (async, escalation tier) · rapidfuzz · Gemini Flash → Groq GPT-OSS 120B → DeepSeek (REST; models verified live against each provider catalog, Aug 2026) · JSONL → PostgreSQL/Neo4j/pgvector · Google Sheets export · pytest + GitHub Actions CI.

0. System Overview



Component	Implementation (trial)	Production swap	Scaling knob
URL frontier	asyncio.Queue per vertical	Redis Streams / SQS	queue depth, consumer count
Crawl workers	coroutines, global semaphore 32, per-domain 4	container replicas	MAX_CONCURRENCY, replica count
Seen-set (dedup)	append-only file, atomic claim()	Redis SET NX + TTL lease	Redis cluster shards
LLM extraction	3-tier fallback chain, per-provider token bucket	same code, more keys/tiers	RPM per tier, tier count
Validation	schema check on every record pre-write	unchanged	—
Entity resolution	seed DB + rapidfuzz, decision log	+ pgvector embedding assist	seed DB size, threshold
Sink	append-only JSONL	Postgres COPY batches	batch size, partitions

Every stage is **stateless against its queue** and **idempotent** (writes keyed by URL fingerprint), so the recovery strategy for any failure is uniformly *retry the unit*, and horizontal scale is purely an infrastructure change.

1. Scale Strategy — 500,000+ records without manual intervention

Discovery is decoupled from extraction. Sitemap/index/API-pagination crawlers only *emit* URLs into the frontier; detail workers only *consume*. Neither knows the other's speed — backpressure comes from bounded queues, keeping memory flat at any frontier size.

Prefer paginated APIs, fall back to sitemaps, HTML last. arXiv API (~50k records/vertical/day within politeness limits), Hugging Face Papers API, the YC company dataset, Greenhouse/Lever boards — all cursor-resumable, so a crashed worker restarts from its checkpoint, not from zero.

Capacity math (why 500k is an infra knob, not a rewrite). One async worker at per-domain politeness 4 req/s sustained ≈ 14k pages/hour/domain. 500k records spread over ~40 source domains ≈ 12.5k/domain — **under one hour of crawling on a single 8-vCPU node** if sources allowed it; in practice politeness and anti-bot pacing dominate, so the planning number is 24–48h with 3 nodes. LLM side: ~2k input tokens/record → 1B tokens for 500k records; at three tiers × published TPM this is throughput-bounded, so the orchestrator treats tiers as *parallel capacity* (round-robin under load) rather than strictly serial fallback.

Unattended operation = idempotency + checkpoints + the failure-mode matrix (§2.3). There is no failure class whose response is "page a human": every class maps to retry, backoff-and-retry, escalate-fetch-tier, or skip-and-log.

2. Handling 413s & 429s across thousands of concurrent extractions

2.1 — 413 / context overflow: never send an oversized payload

- Reduce:** HTML → dense text via lxml (scripts/nav/boilerplate stripped) — typically a 10–20× byte reduction before any budgeting.
- Estimate client-side:** chars/4 heuristic with a 0.9 safety margin against the *per-tier* input budget (Gemini 100k tokens, Groq 16k, DeepSeek 48k) minus prompt overhead.
- Truncate head+tail-biased** (70/30): titles, meta, and lead paragraphs carry the signal in news/job pages; the middle is dropped first, marked ...[truncated]... so the model knows.

4. **Belt-and-braces:** if a provider still returns 413 (estimator too optimistic for its tokenizer), the payload is **halved and retried** — geometric convergence guarantees success in ≤ 3 halvings.

2.2 — 429 / rate limits: proactive pacing + reactive backoff

- **Proactive:** a token-bucket limiter per provider (aiolimiter, keyed to published RPM) smooths thousands of concurrent tasks *before* the API sees them. This is the difference between "we handle 429s" and "we rarely cause them".
- **Reactive:** exponential backoff with **full jitter** — `sleep = uniform(0, min(60s, 1s·2^attempt))` — honoring `Retry-After` (numeric or HTTP-date). Full jitter specifically prevents synchronized retry stampedes when a whole worker fleet gets limited at the same instant.
- **Fallback chain as load-shedding:** ≥ 3 consecutive 429s on a tier routes traffic to the next tier instead of queueing — the chain is simultaneously a *reliability* mechanism (provider outage), a *correctness* mechanism (unparseable JSON → next tier), and a *throughput* mechanism (rate pressure). The winning provider is stamped on every record (`extractedBy`) for audit.

2.3 — Failure-mode matrix (crawl + LLM)

Signal	Diagnosis	Automated response
HTTP 429 / <code>Retry-After</code>	rate limited	full-jitter backoff, honor header, ≤ 3 attempts → next tier/proxy
HTTP 413 / context error	payload too large	halve payload, retry (≤ 3 halvings)
HTTP 403/503 + challenge page	anti-bot block	rotate UA → rotate proxy → escalate to Playwright tier
HTTP 404/410	gone	skip, log, never retry
5xx / timeout / conn reset	transient	jittered retry ≤ 3 , then requeue unit
LLM output unparseable	model drift	1 retry same tier → next tier → deterministic heuristics
All LLM tiers exhausted	keys/quota	record persists un-enriched; idempotent <code>enrich</code> backfills later

3. Freshness Tracking — never process the same item twice, across nodes

- **Canonical URL fingerprint:** lowercase scheme+host, tracking params (`utm_*` , `fbclid` , ...) stripped, query keys sorted, trailing slash collapsed → SHA-256. This fingerprint is the *global* dedup key for both bulk and fresh verticals.
- **Atomic claim semantics:** the trial store is an append-only local file with `claim(fp) → bool`; the production store is Redis `SET fp NX EX <lease>` — the same one-method interface, so distribution is a backend swap. `NX` makes check-and-claim atomic: two nodes can never both win a URL. The TTL lease means a worker that crashes *after claiming but before writing* releases the URL automatically.
- **Two-phase freshness gate, before any LLM spend:** (1) RSS/API dates filter stale items pre-fetch; (2) the article page itself is re-verified — meta tags (`article:published_time`), JSON-LD `datePublished` , `<time datetime>` , then visible relative dates ("2 hours ago") normalized to UTC. Items that cannot *prove* ≤ 24 h freshness are **dropped, never guessed** — each kept record carries `dateConfidence` (`meta` | `json_ld` | `time_tag` | `relative_text` | `rss`).
- **No-date heuristic:** date-less sources fall back to first-seen semantics on a **content hash** of extracted text (whitespace-normalized), so republished-URL tricks don't double-ingest and layout changes don't re-trigger.
- **Fuzzy-date safety:** free-text date parsing is only trusted when the text visibly contains a year or month name — otherwise parsers hallucinate dates out of strings like "Page 20 of 30" and silently poison the freshness guarantee.

4. Storage Strategy

Layer	Choice	Justification
System of record	PostgreSQL (JSONB)	Versioned envelope fits JSONB with real indexing: <code>UNIQUE(url_fingerprint)</code> upserts, GIN on <code>content->>'entityName'</code> , B-tree on <code>collectedAt</code> ; transactional; trivially handles 500k–50M rows on one primary; <code>schemaVersion</code> enables in-place migrations.
Ingest buffer	Append-only JSONL	Crash-safe (a torn line is discardable, never corrupting), resumable, streams into Postgres <code>COPY</code> and the Sheets exporter; what the trial ships.
Relationship graph	Neo4j	The product is an intelligence graph. Edges: <code>(Startup)-[:BUILDS]->(Product)</code> , <code>(Paper)-[:IMPLEMENTED_BY]->(Repo)</code> , <code>(Startup)-[:HIRING]->(Job)</code> , <code>(News)-[:MENTIONS]->(Startup)</code> . Multi-hop questions ("startups whose papers trend on GitHub <i>and</i> are hiring researchers") are 3-join-hostile in SQL, native in Cypher.

Layer	Choice	Justification
Similarity assist	pgvector (inside Postgres)	Embedding search for near-duplicate entities and article dedup beyond exact hashing; kept in-Postgres to avoid operating a second database before scale demands it.

One source of truth: JSONL → Postgres; Neo4j and pgvector are *projections* rebuilt from Postgres at any time — no dual-write consistency problem. Entity resolution runs pre-projection so the graph only ever contains canonical nodes; the raw→canonical mapping log is itself a table (provenance + resolver regression testing).

5. Operations, Observability, Integrity

- **Metrics per vertical:** records written/dropped-invalid/dropped-stale, dedup hit-rate, fetch success by tier (plain/proxy/rendered), LLM calls by provider × outcome, 429/413 counts, backfill gap size. Exported as counters; alert on *rates* (e.g., >20% stale-drop spike = a source changed its date markup).
- **Structured logging:** namespaced (`frontieratlas.crawlers.news`), timestamped, levels env-controlled; every drop states its reason — nothing disappears silently.
- **Security:** secrets only via `.env/environment` (gitignored; `.env.example` documents shape); no keys in code or logs; robots-aware politeness caps per domain.
- **CI:** GitHub Actions runs compile + 21 unit tests (dates, dedup, resolution, chunking, backoff, schema validation) on every push.
- **Data integrity (anti-hallucination):** every record must carry a live `source.url`; LLMs only ever *restructure fetched text* — extraction prompts forbid facts not present in the payload, and records failing schema validation are dropped and counted, not repaired by guesswork.